

ADIOS at scale

Presenter: Norbert Podhorszki

Oak Ridge National Laboratory, Computer Science and Mathematics Division

Excalibur - US extreme data workshop and tutorials

Cambridge, UK

July 21, 2023

I/O Challenges

- Write/read a TB dataset frequently enough for science
- Write a 60 TB checkpoint every half an hour (particle codes)
- Write/Read data from $O(100k)$ processes (CPU-only codes, Arm systems)
- Collect data from thousands of applications to train AI/ML (DeepDriveMD)
- Couple two/more parallel applications in every iteration (Fusion Whole Device Modeling)
- Write a few PBs of data produced by an ensemble occupying the entire machine (Earth-scale tomography)
- Staging data processing pipeline that can process data TB/s (SKA epoch pipeline on Summit)
- Replace costly parts of simulation with a surrogate model tightly coupled to the simulation using staging (correlated electrons simulation, WRF, fusion, particle accelerators)
- Scalable I/O library in number of processors, number of variables, number of steps, frequency of steps

ADIOS: high-performance publisher/subscriber I/O framework:

Approach

- Easy-to-use, high performance I/O **abstraction** to allow for on-line/off-line memory/file data subscription service
- Sustainable solution that works with multi-tier storage and memory systems for self-describing data-streams

Details

- Declarative, publish/subscribe API separated from the I/O strategy
- Multiple implementations (engines) provide functionality and performance
- Rigorous testing ensures portability
- GPU-aware to reduce data movement
- <https://github.com/ornladios/ADIOS2>

API:

- Each process describes what data it has/needs
- Forced to declare the I/O phase

I/O strategy:

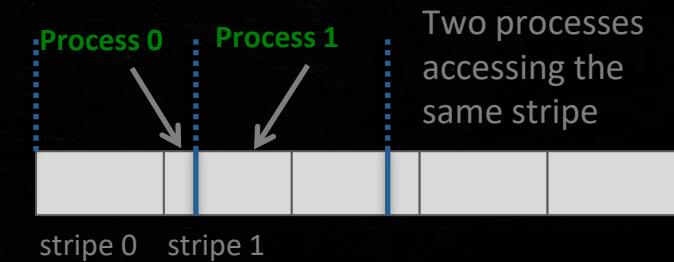
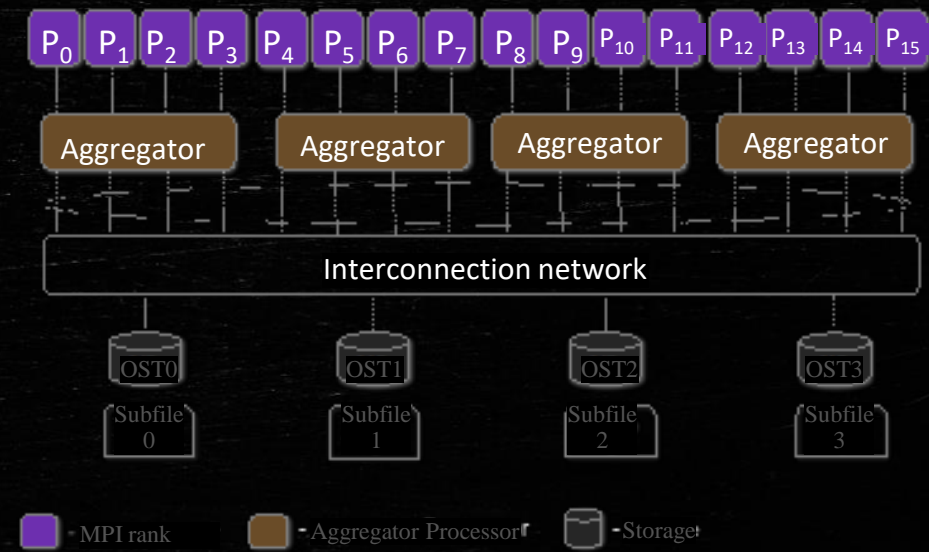
Media, aggregation, number of sub-files, target file-system hacks, file format not expressed at the code level



File format + API designed for scalability

- Avoid latency (of small writes): **Buffer** data for large bursts – use a type of self-describing log file format
- Avoid accessing a file system target from many processes at once
 - **Aggregate** to a small number of actual writers:
 - Avoid lock contention
 - **Striping correctly & writing to subfiles**
- Avoid global communication wherever possible

Liu, Q., Klasky, S., et al. "Hello ADIOS: the challenges and lessons of developing leadership class I/O frameworks." *Concurrency and Computation: Practice and Experience* 26.7 (2014): 1453-1473.



Application	Nodes/GPUs	Data Size/step	I/O speed
SPECFEM3D	3200/19200	250 TB	~2 TB/sec
GTC	512/3072	2.6 TB	~2 TB/sec
XGC	512/3072	64 TB	1.2 TB/sec
LAMMPS	512/3072	457 GB	1 TB/sec

BP File format

- Dataset is a **directory**, not a single file
 - Data is stored in multiple files
- Arrays (variables) are stored as **chunks**, that can be **compressed**/decompressed individually
 - One arrays spans all files, multiple arrays (chunks) are stored in each file
 - Don't confuse this with HDF5/Zarr's static chunk concept
- **Metadata** is stored **separately** from data (separate files)
- Metadata is a binary format and is big
 - information per chunk has to be stored to be able to find, locate and retrieve it

Why use ADIOS over other technologies

- **Scalable** performance
 - designed to be a parallel IO library,
 - one application can utilize the bandwidth of the entire file system (when running at large scale)
 - proven by users for many PBs datasets, tens of thousands of nodes, daylong runs, thousands of steps
- Scalable (distributed) **compression** routines
 - lossless and lossy, blosc, bzip2, zfp, sz, mgard
- Safe to **append** data into single dataset during run, at scale
 - new writes cannot corrupt existing data
 - entire output “step” is either in or not
 - application failure will not corrupt previously written data/file
- Able to **read** data from files **being written**
 - previous steps only, not the one under construction
 - stable semantics: data in a published step is always completely available

ADIOS Scaling for large parallel file systems

- There are a **few** options for changing the performance of ADIOS for your code on all HPC systems
 - Aggregation – choosing the number of sub-files for maximizing the filesystem bandwidth
 - Appending multiple steps into a single output for minimizing the filesystem metadata overhead
 - Asynchronous write with BP5 engine
 - Choosing the “best” ADIOS engine/storage system (staging, NVRAM-flush) for minimizing the variability

ADIOS Scaling for large parallel file systems: number of files


- Not good:
 - Single, global output file from many writers (or for many readers)
 - Bottleneck at file access
 - One file per process
 - Chokes the metadata server of the file system at create
 - Reading from different number of processes is very difficult
- Good:
 - Create K subfiles where K is proportioned to the capability of the file system, not the number of writers/readers
- ADIOS BP engine options
 - **NumAggregators**
 - **AggregatorRatio**
- ADIOS default settings
 - One file per compute node

```
<io name="SimulationOutput">  
  <engine type="BP5">  
    <parameter key="NumAggregators" value="2048"/>  
  </engine>  
</io>
```


Example: VPIC I/O test on Summit

- A fixed aggregation ratio breaks down as we scale up the nodes
- Best options here:
 - 42 nodes – $42 \times 42 = 1764$ subfiles (1:1)
 - 84 nodes – $84 \times 42 = 3528$ subfiles (1:1)
 - 168 nodes – $168 \times 21 = 3528$ subfiles (1:2)
 - 336 nodes – $336 \times 6 = 2016$ subfiles (1:7)
- Summit general guidance
 - One subfile per GPU (6 per node) is a good start but apps usually have one MPI task/GPU, so keep the total number of subfiles below 4000


chart courtesy of Junmin Gu, LBNL



Application	Nodes/GPUs, 6 tasks/node	Data Size per step	I/O speed	ADIOS NumAggregators
SPECFEM3D_GLOBE	3200/19200	250 TB	~2 TB/sec	3200 (1:6 aggregation ratio)
GTC	512/3072	2.6 TB	~2 TB/sec	3072 (1:1 aggregation ratio)
XGC	512/3072	64 TB	1.2 TB/sec	1024 (1:3 aggregation ratio)
LAMMPS	512/3072	457 GB	1 TB/sec	512 (1:6 aggregation ratio)

ADIOS Scaling for large parallel file systems: number of steps

- Another aspect of number of files: number of output steps
- New output every step -> many files over the course of simulation
 - If the rate of steps stresses the file system, write performance will drop
 - Actually, the total time of creation of files will add up
- Appending multiple output steps into same file is better





XGC 100 output steps in 1245 second simulation,
40 GB per step (4TB total)

- One file per node created in each step
- Single output for all steps (one file per node)

Optimized GTC, a fusion PIC code, I/O on Summit

PI: Zhihong Lin, UC Irvine

- Change to ADIOS I/O: Total simulation time reduced from 9.5 hours to 6.1 hours on 1024 nodes on Summit




*** average cost when measuring 10 000 steps**

WarpX example: appending + aggregation


PI: Jean-Luc Vay, LBNL

- One file per compute node (6 processes on Summit)
- Better performance at 512 nodes and over
- One file per rank at smaller jobs

WarpX on Summit, Original vs
ADIOS new output per step vs
ADIOS single output




WarpX on Summit, 6 rank per node setup
240 TB on 1024 nodes (6144 cores)



240 TB on 1024 nodes (6144 cores)


Single process I/O examples: NWChem, LAMMPS

- **Moves data** between processes as part of preparation for I/O
 - “I am doing POSIX/Fortran I/O on rank 0”
 - While gathering data, no one is writing
- **Single file** output – not utilizing available bandwidth



ADIOS can write all steps out with little cost
(here every 0.2 seconds)

Summit 512 nodes
12B atoms, 5 TB





<https://github.com/lammps/lammps/tree/master/src/USER-ADIOS>

- USER-ADIOS package in LAMMPS for dump commands
 - dump atom/adios
 - dump custom/adios
- Output goes into an I/O stream
 - BP4 file by default
 - Can use staging engines

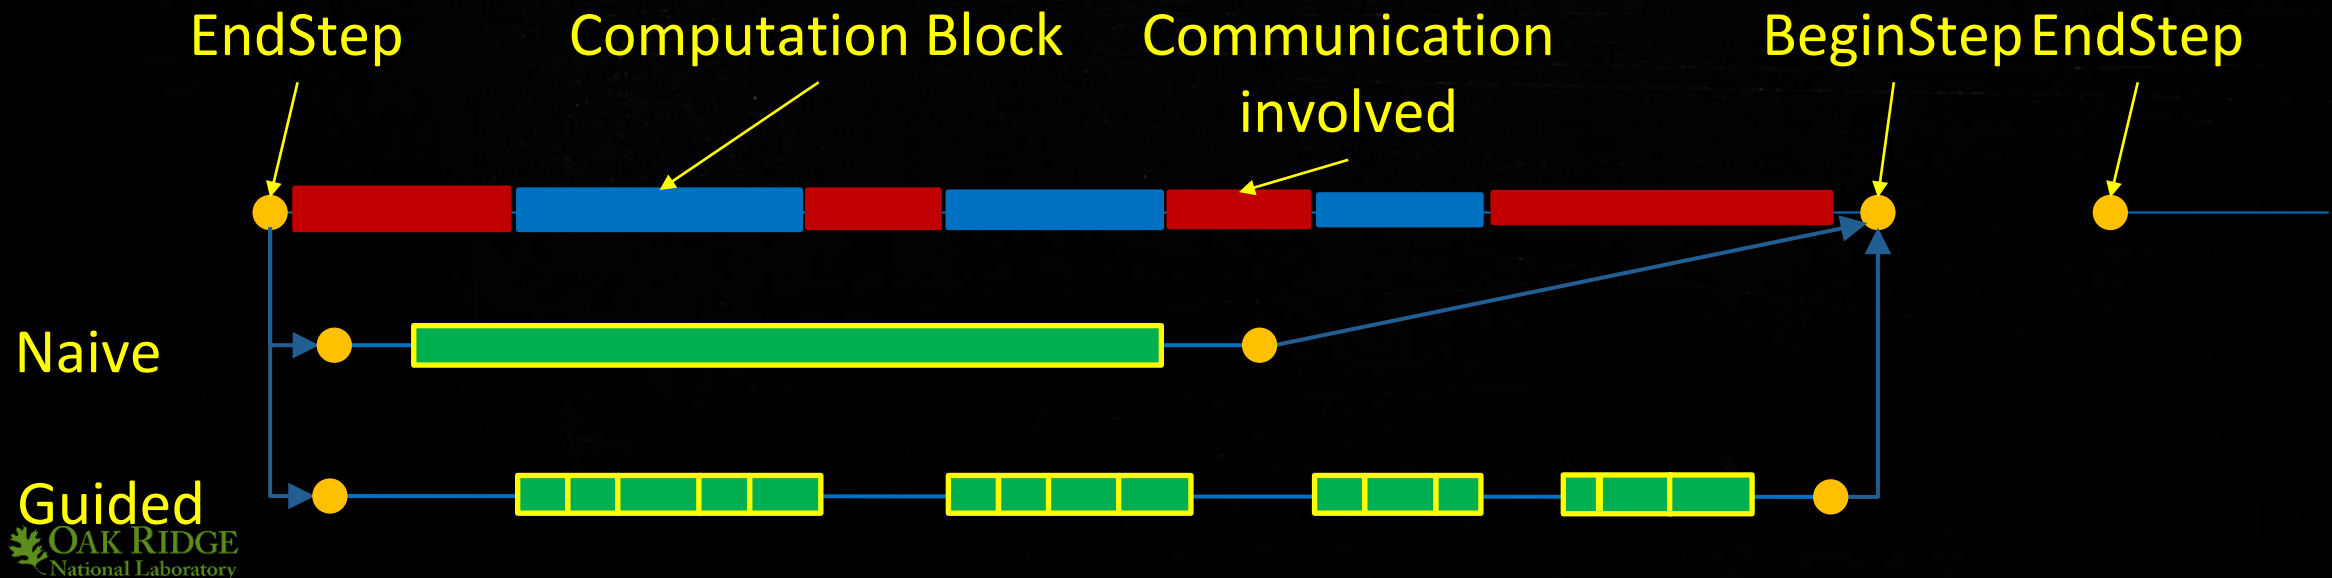
Asynchronous write to storage (BP5 engine, 2.8.0 release)

- User friendly On/Off option
 - No need to modify the user code
- Only data writing is async, metadata gathering and writing is still sync





Async strategies

- Naive
 - dump data without thinking
- Guided
 - Application augmented with `EnterComputationBlock/ExitComputationBlock` pairs
 - Attempt writing during computation blocks
 - Naïve dump when running out of (forecasted) computation blocks



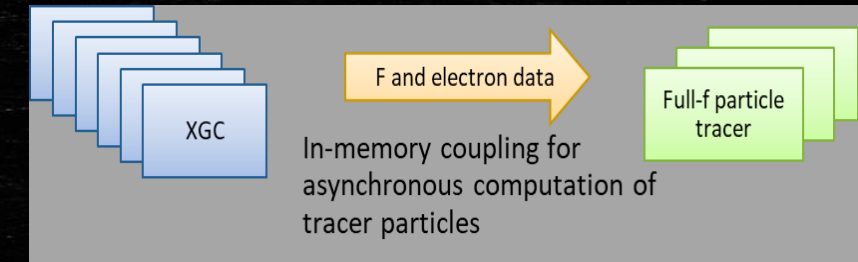
Async IO with XGC only small data

- This is a small D3D run case and quick test



Staging use case: off-load non-scaling code part

- Tracer particle analysis enables understanding of the transport characteristics spanning the pedestal and scrape-off layer
- It is costly to perform and is communication-heavy
- Asynchronously stage data to the tracer particle analysis running on additional nodes
 - Coupled data size: $f_0(95 \text{ GB}) + E_{\text{rho}}/\text{pot_rho}(1.4\text{GB})$
- Reduced 36% of the XGC iteration time by using asynchronous services (only 0.4% time-overhead for coupling data)



Full-f particle analysis coupling



Full-f coupling performance on Summit with ADIOS using 4/1024 extra nodes

Choi, J. Y., Chang, C. S., Dominski, J., Klasky, Churchill, M., S., Merlo, G., Suchyta, E., et al.
“Coupling exascale multiphysics applications: Methods and lessons learned”. IEEE e-Science, 2018.

Summary

ADIOS brings a programming interface and a framework of many solutions to the generic problem of producing and consuming data

- The interface frees scientists from the limited scope of file-based data processing
 - Being fully applicable to file-based data processing
- Offering a bridge from their scientific workflows that work now to the future, where they will extend their workflows with
 - More efficient data processing
 - Interactive visualization
 - Code coupling
 - On-the-fly AI training
 - Combining experimental data with simulation data